

Summary

AI Engineer experienced in designing, fine-tuning, and deploying scalable ML and LLM-powered systems. Proficient in building robust AI pipelines, integrating MLOps workflows, and leveraging tools like LangChain, Hugging Face, and Streamlit for NLP, generative, and predictive applications. Proven ability to deliver intelligent solutions that automate processes, extract insights, and solve real-world problems.

Skills

Programming Languages & Tools: Python, SQL, Git

AI & ML: Supervised/Unsupervised Learning, Generative AI, Agentic AI, LLM Fine-Tuning, Prompt Engineering

Frameworks & Libraries: Scikit-Learn, TensorFlow, PyTorch, LangChain, LangGraph, LangSmith

MLOps & Deployment: MLflow, FastAPI, Streamlit, Docker, CI/CD Pipelines

Cloud & Infra: Microsoft Azure, Amazon Web Service (AWS)

Databases & Vector Stores: PostgreSQL, MySQL, MongoDB, FAISS, ChromaDB

Automation & Workflow Tools: Vapi, Make, n8n

Visualization: Power BI, Matplotlib, Seaborn, Plotly

Languages: English (Professional), Urdu (Native)

Experience

AI Developer

Bright Solutions | August 2025 - Present | Hybrid

- Delivering end-to-end client AI products with a focus on automation, Generative AI, and agentic AI solutions.

AI Intern

SaylaniTech Limited | July 2025 - Present | Onsite

- Contributing to LLMs, NLP, and AI product development with modular pipelines, deployment, and API integration.

AI Intern

Visdalytics | May 2025 | Remote (1-month)

- Built AI chatbots, TTS tools, Roman-to-Urdu converter, and CV/ML pipelines (YOLOv8, Transformers, TensorFlow).

Projects

• DineMate: Agentic AI for Automated Food Ordering

Personal Project | 2025

- **Challenge:** Replaced manual food ordering with a **real-time**, multi-role AI system featuring voice input and analytics.
- **Action:**
 - Engineered **LangGraph-based multi-agent** backend to **orchestrate** orders, kitchen, and admin tasks.
 - Developed a **voice-enabled** food ordering **chatbot** using Qwen, Whisper ASR, Torch TTS, LangChain, and SQLite.
 - Integrated **SQLite** to simplify **local deployment** and implemented **secure, role-based authentication**.
 - Designed **analytics dashboards** to monitor **peak hours**, **sales trends**, and **customer spend patterns**.
- **Result:**
 - Boosted **ordering efficiency** by **30%**, reducing **manual operations** across all user roles.
 - Enhanced **decision-making** via **dashboards**, improving **business visibility** by **25%**. [Live Demo](#).

• SupportGenie: Dual Fine-Tuned LLM Customer Support Chatbot

Personal Project | 2025

- **Challenge:** Upgraded a generic **RAG** bot into a **domain-specialized assistant** using **fine-tuned LLMs** on banking FAQs.
- **Action:**
 - Fine-tuned **Mistral-7B** and **LLaMA-3 (8B)** using **QLoRA** on a custom Hugging Face-hosted [dataset](#).
 - Built a Streamlit app with **MongoDB**, **FAISS** retrieval, **dual-LLM inference**, and **sentiment analytics**.
 - Benchmarked model outputs, selecting **Mistral** for its **domain alignment** and concise **factual responses**.
 - Released a Colab **notebook** and integrated real-time usage **dashboards** with sentiment tracking.
- **Result:**
 - Improved answer relevance by **20%** versus **base models** through targeted fine-tuning and **validation**.
 - Deployed a **public chatbot** with **dual inference**, **feedback analytics**, and FAQ explainability. [Live demo](#).

• **Diagnosify: LLM-Powered Medical Report Insights**

Hackathon Project | 2025

- **Challenge:** Engineered a reliable AI system to interpret and explain unstructured medical reports using LLMs.
- **Action:**
 - Combined OCR, prompt-tuned LLMs, and PDF parsing to extract structured clinical data.
 - Built a RAG-enabled chatbot for real-time, document-grounded medical question answering.
 - Integrated RAGAS to evaluate answer faithfulness and ensure response accuracy.
 - Logged all chatbot interactions and metrics in MongoDB for auditability and analysis.
- **Result:**
 - Delivered a production-ready health AI tool within 12 hours during live hackathon.
 - Demonstrated trustworthy, domain-specific LLM performance with built-in evaluation logging. [Live demo](#).

• **ChurnSage: MLOps-Driven Customer Retention System**

Personal Project | 2025

- **Challenge:** Predicted customer churn from imbalanced telecom data with scalable ML infrastructure.
- **Action:**
 - Built a full ML pipeline with SMOTEENN, encoding, scaling, and hyperparameter tuning using GridSearchCV.
 - Trained and tracked Logistic, Random Forest, and SVC models using MLflow for reproducibility.
 - Leveraged Azure Cloud Notebooks for scalable training and managed experiment logging.
 - Deployed the best model with FastAPI (REST API + SQLite) and integrated LLMs to explain user churn from reviews.
- **Result:**
 - Achieved ~94% test accuracy with robust handling of class imbalance and optimized modeling.
 - Delivered a CI-ready, batch-enabled ML pipeline with LLM explainability. [Live demo](#).

Certifications

- Deep Learning Specialization - [DeepLearning.AI](#) (Coursera), 2025
- Natural Language Processing Specialization - [DeepLearning.AI](#) (Coursera), 2025
- Machine Learning Professional Certificate - [IBM](#) (Coursera), 2024
- Artificial Intelligence and Data Science - [S.M.I.T.](#), 2025

Education

Sindh Madressatul Islam University

Karachi, Pakistan

B.S. in Computer Science

- **Relevant Coursework:** Machine Learning, Deep Learning, NLP, Data Structure & Algorithms, Web Technologies.

Achievements

- Top 10 Finalist - SMIT AI Hackathon (2025) for Diagnosify: Medical Insight System.